

Introduction to Java Pathfinder

Cyrille Artho

KTH Royal Institute of Technology, Stockholm, Sweden
School of Electrical Engineering and Computer Science
Theoretical Computer Science

`artho@kth.se`

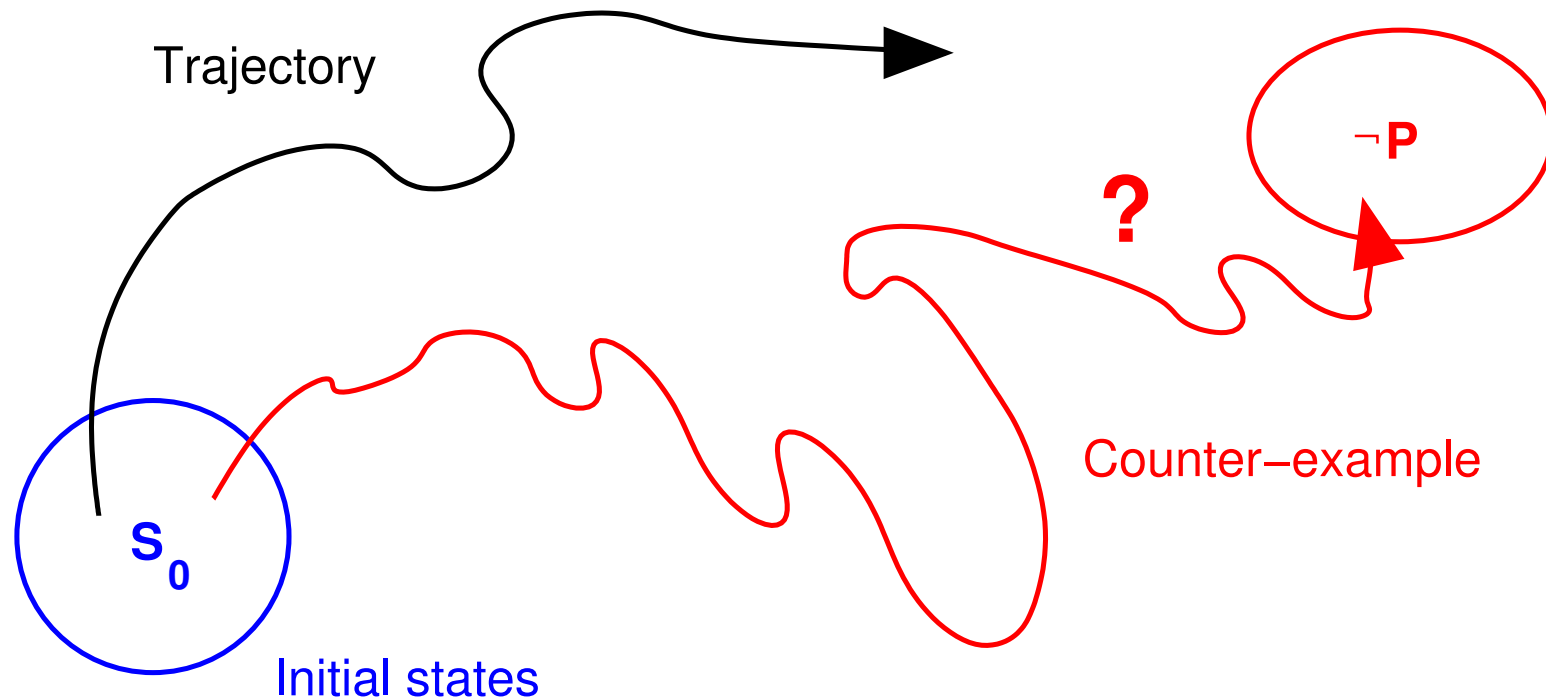
Today's objectives

1. Model Checking vs. Software Model Checking
2. Introduction to Java PathFinder:
 - (a) Basic usage.
 - (b) `fork/join`.
 - (c) Deadlock detection.
 - (d) `wait/notify`.

Understand JavaPathFinder and its output.

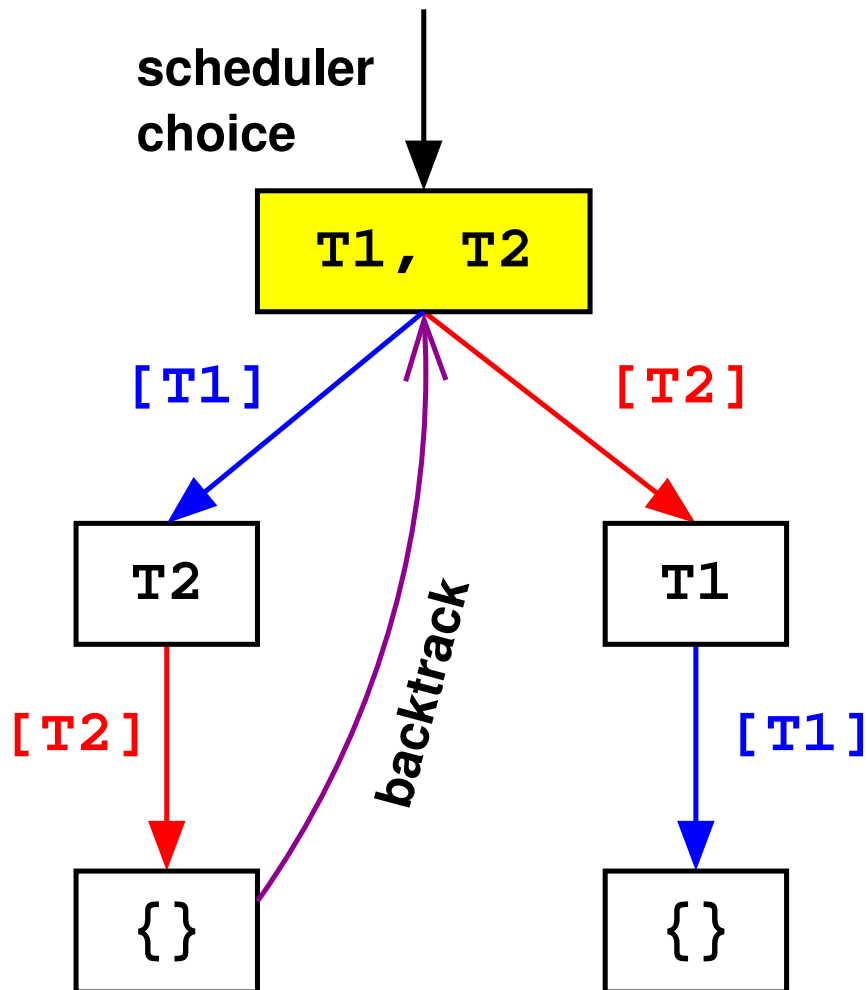
- Furthermore, crash course about thread control and locking in Java.

Model Checking = state space search



- Traditionally applied to specifications, protocols, algorithms.
- Certain types of software (embedded) can be mapped to such model checkers.

Software model checking




- MC can **backtrack** program execution.
- Explores all possible thread schedules!
- Finds all possible program failures.

Key limitation: Scalability.

Another limitation: Networking!

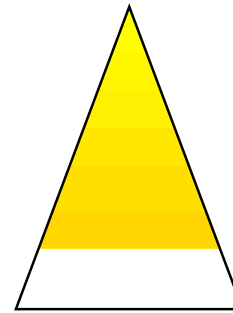
Model checking vs. testing

Testing



- **Only one trace.**
- No rules (but assertions).
- **May miss defects.**
- **Scalable.**

Software Model Checking (SMC)




- Many (all) traces.
- Monitoring rules.
- **Finds all defects.**
- **Resource hungry.**


- SMC is a rigorous formal method.
- SMC does not suffer from false positive (like static analysis).
- SMC provides traces (execution history) when it finds a defect.

Heuristic model checking

Full Verification



- Formal verification.
- Prove properties.
- Include all states.
- **Not very scalable.**



Fault-finding

- Partial verification.
- May miss defects.
- Many, but not all states.
- **May find bugs faster.**

- Heuristics:

- Prefer states that likely result in a failure.
- Helpful with a priori knowledge about defect (deadlock? data race?).
- Heuristics may be supported by static analysis:

Rungta, Mercer: A Meta Heuristic for Effectively Detecting Concurrency Errors. HVC 2008.

What is Java PathFinder?

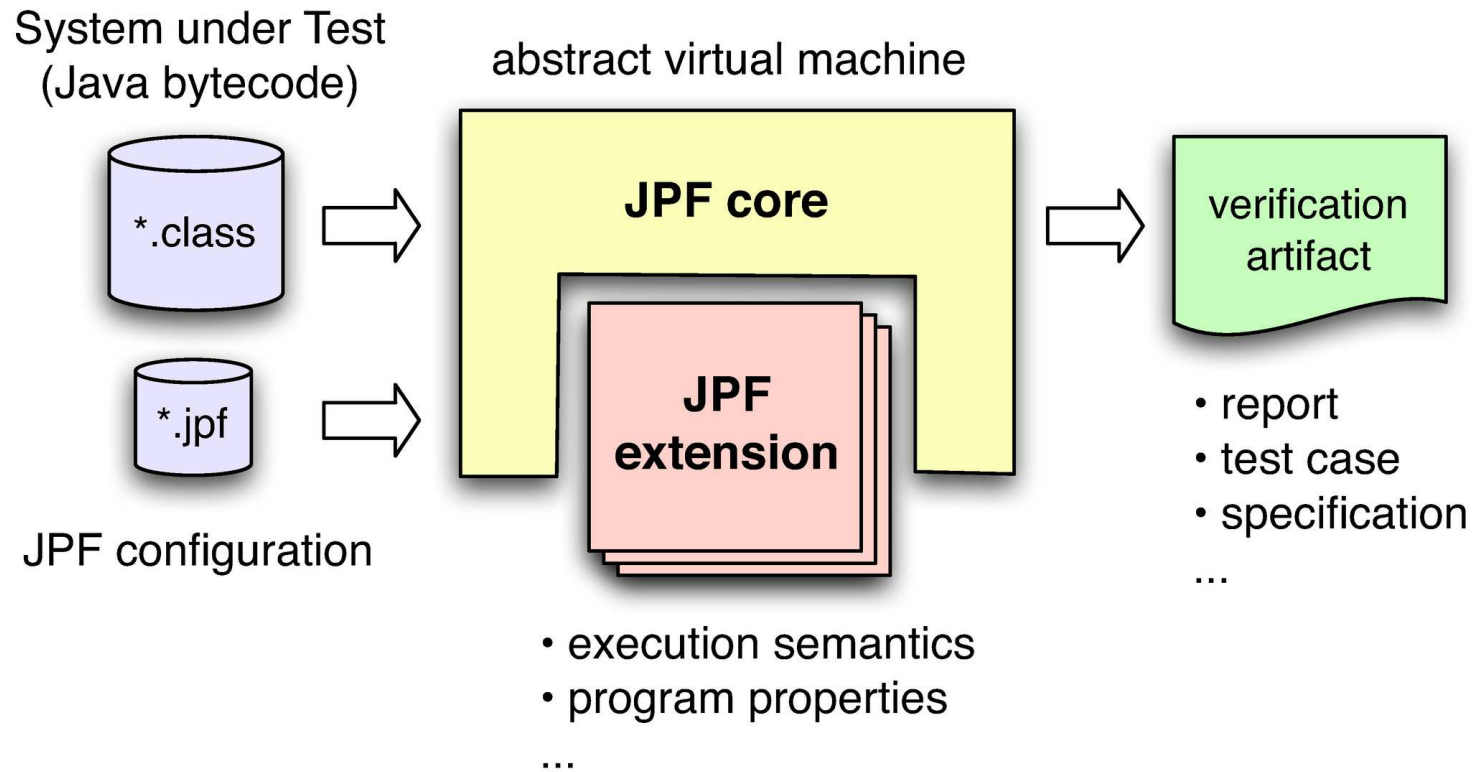


Image source: Java PathFinder tutorial,
<http://babelfish.arc.nasa.gov/trac/jpf/attachment/wiki/presentations/start/jpf-graphics.zip>

- Born as software model checker for Java bytecode.
- Evolved into multi-purpose tool framework.

History of Java Pathfinder

- **1999: Project started** at **NASA Ames** as front end to Spin model checker.
- **2000: Reimplementation** as concrete virtual machine for software model checking (concurrency defects).
- **2003:** Introduction of **extension** interfaces.
- **2005: Open sourced** on sourceforge.
- 2008: First participation in Google Summer of Code.
- 2009: Moved to own server, hosting extension projects and wiki.
- **2017: Moved to github.**
- Many collaborators/users:
 - Academic research: > 20 univ. (USA, Canada, Japan, South Africa, Europe).
 - Companies: Fujitsu, other Fortune 500 companies.

Installing Java PathFinder

- You first need to install **Java 8**, **gradle**, and **git**.
- Clone the JPF source code repository:
<https://github.com/javapathfinder/jpf-core/>
- Run the following commands:

```
mkdir ~/jpf  
cd ~/jpf  
git clone https://github.com/javapathfinder/jpf-core.git  
cd jpf-core  
./gradlew
```
- Configuration files have been prepared to be used with JPF 8.

Testing your JPF installation

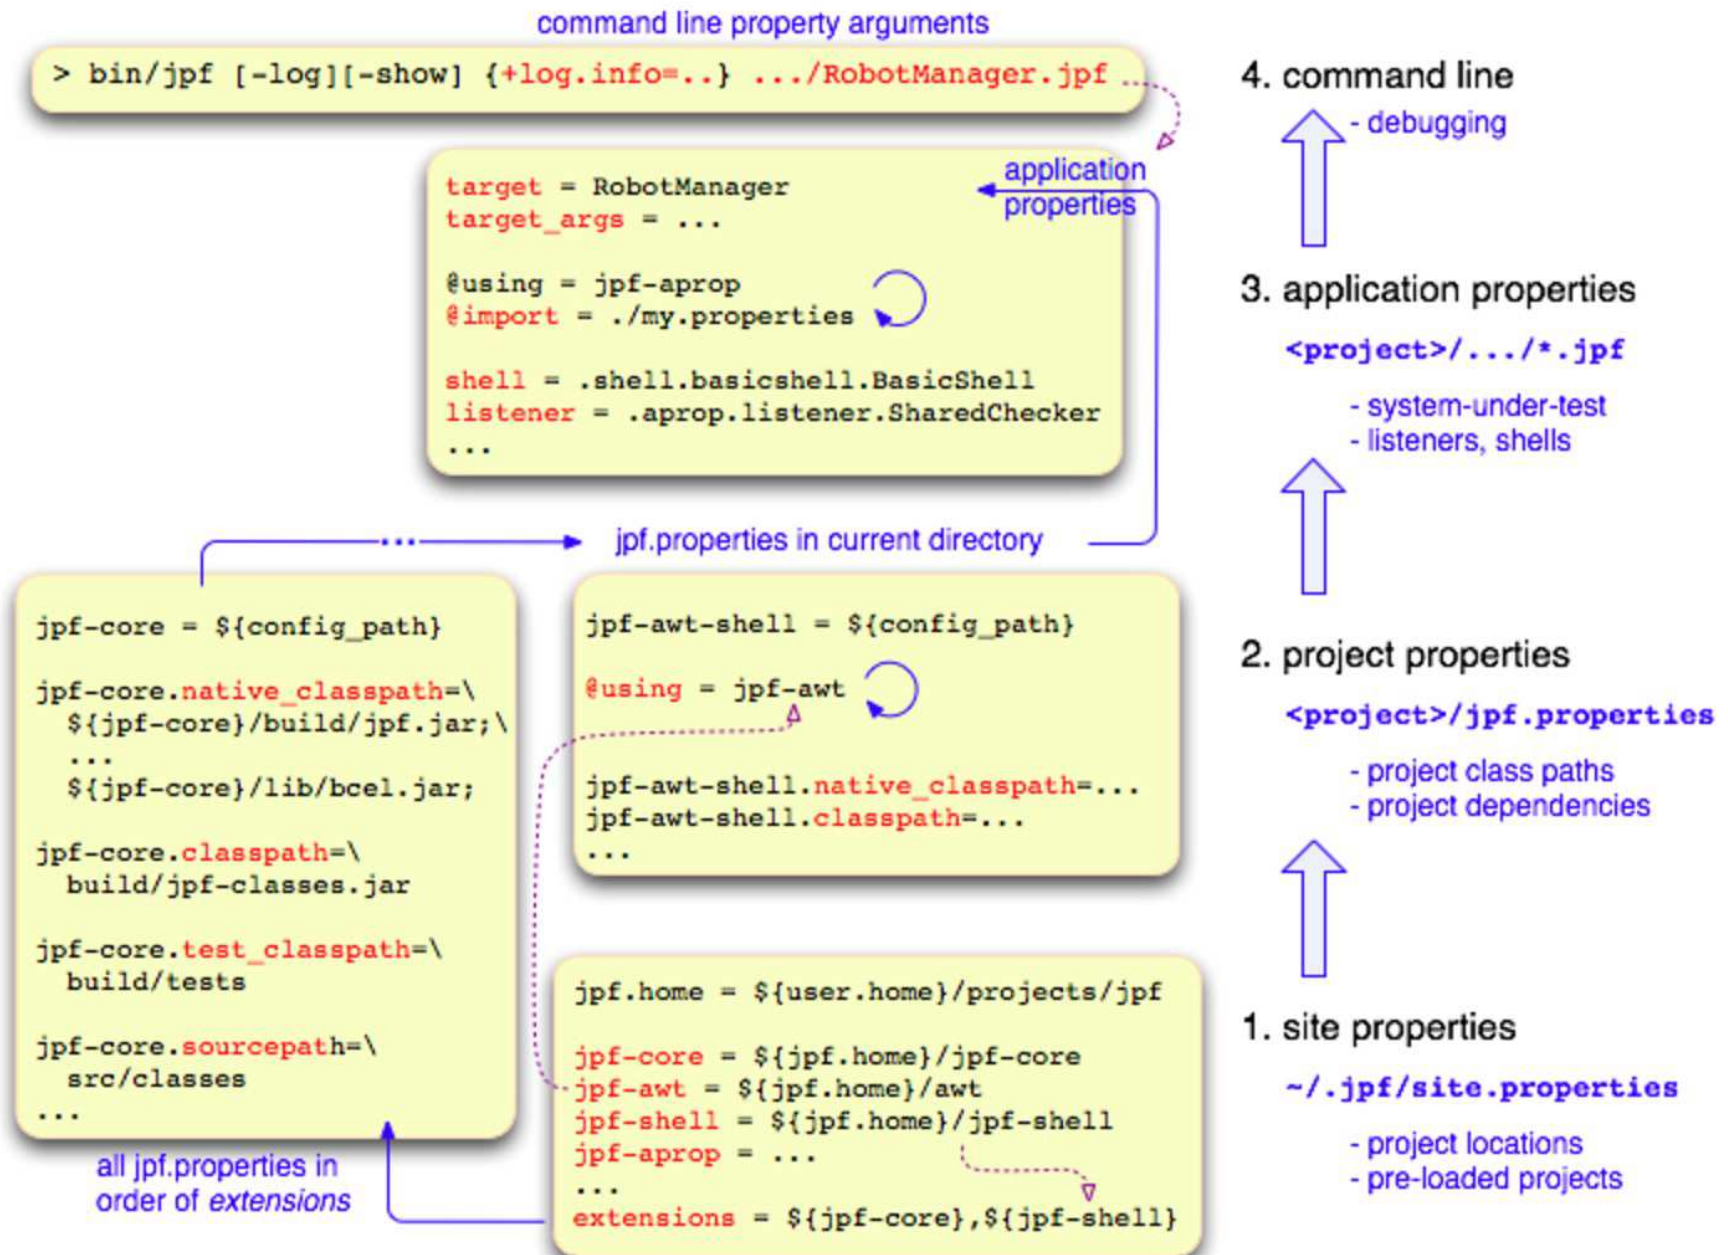
Recommended: Try this before Easter.

The following commands should provide meaningful output:

```
bin/jpf src/examples/HelloWorld.jpf  
bin/jpf -version
```

- Tested on Mac OS X and the lab computer's Ubuntu.
- Second command also prints an error message (no file given); you can ignore that as long as you do not see only an error message.

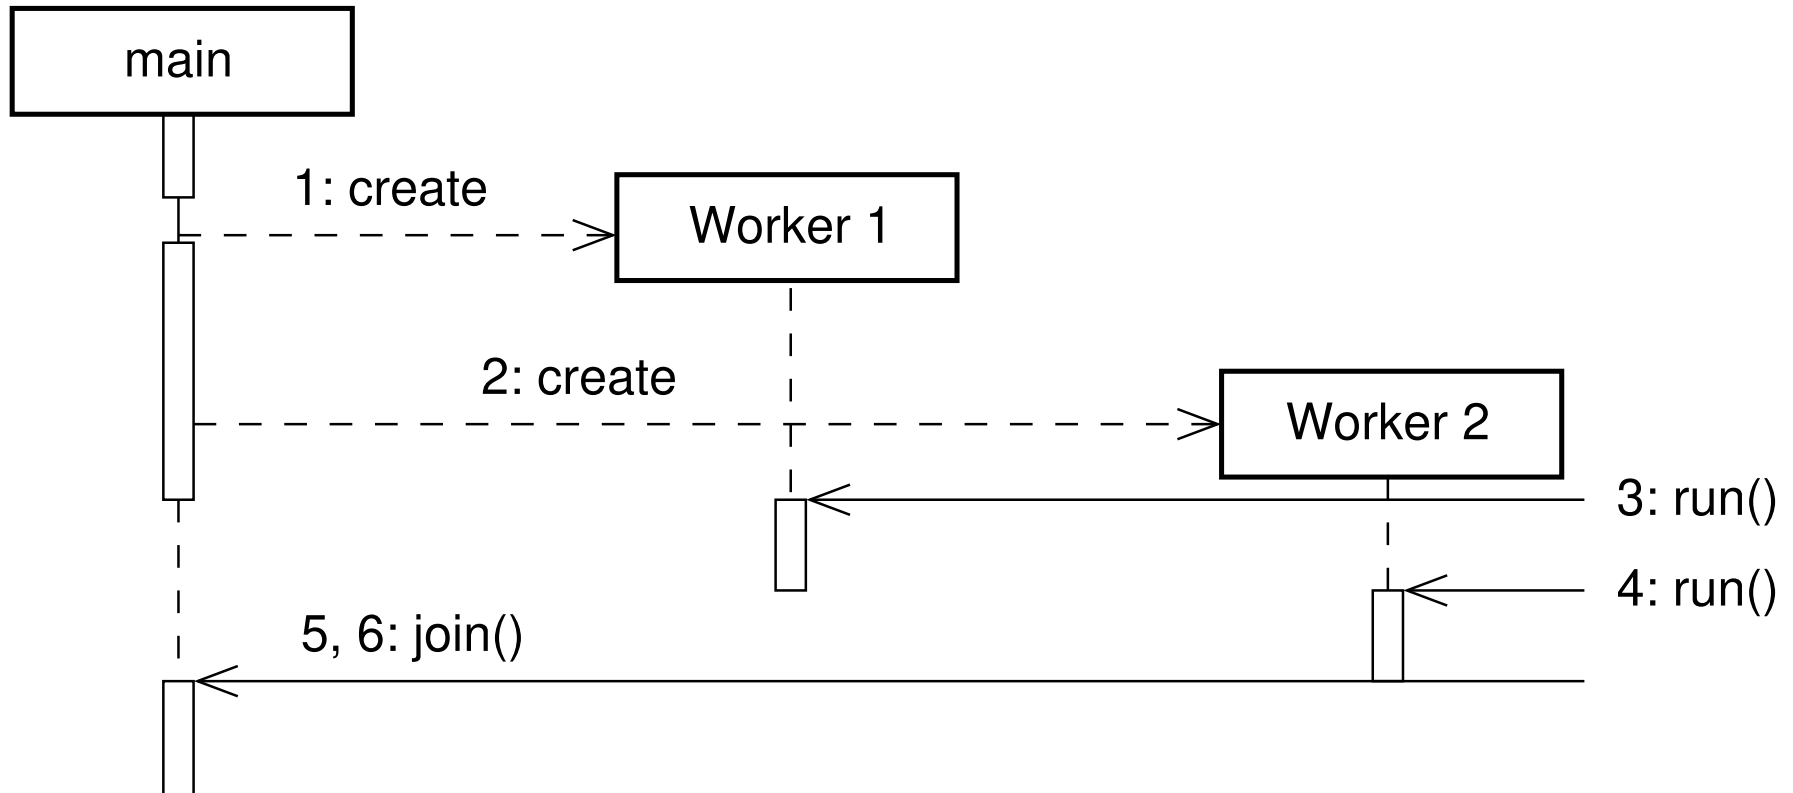
Configuration under JPF 8 (slide from JPF manual)



Example 1: Concurrent „Hello, World!”

```
public class HelloWorld extends Thread {
    public static final void main(String[] args) {
        StringBuffer buf = new StringBuffer();
        HelloWorld hw1 = new HelloWorld(buf, "Hello, ");
        HelloWorld hw2 = new HelloWorld(buf, "World!");
        hw1.start(); // spin off first worker thread
        hw2.start(); // spin off second worker thread
        try {
            hw1.join(); // wait for first worker to finish
            hw2.join(); // wait for second worker
        } catch (InterruptedException e) {
        }
        System.out.println(buf.toString());
    }
}
```

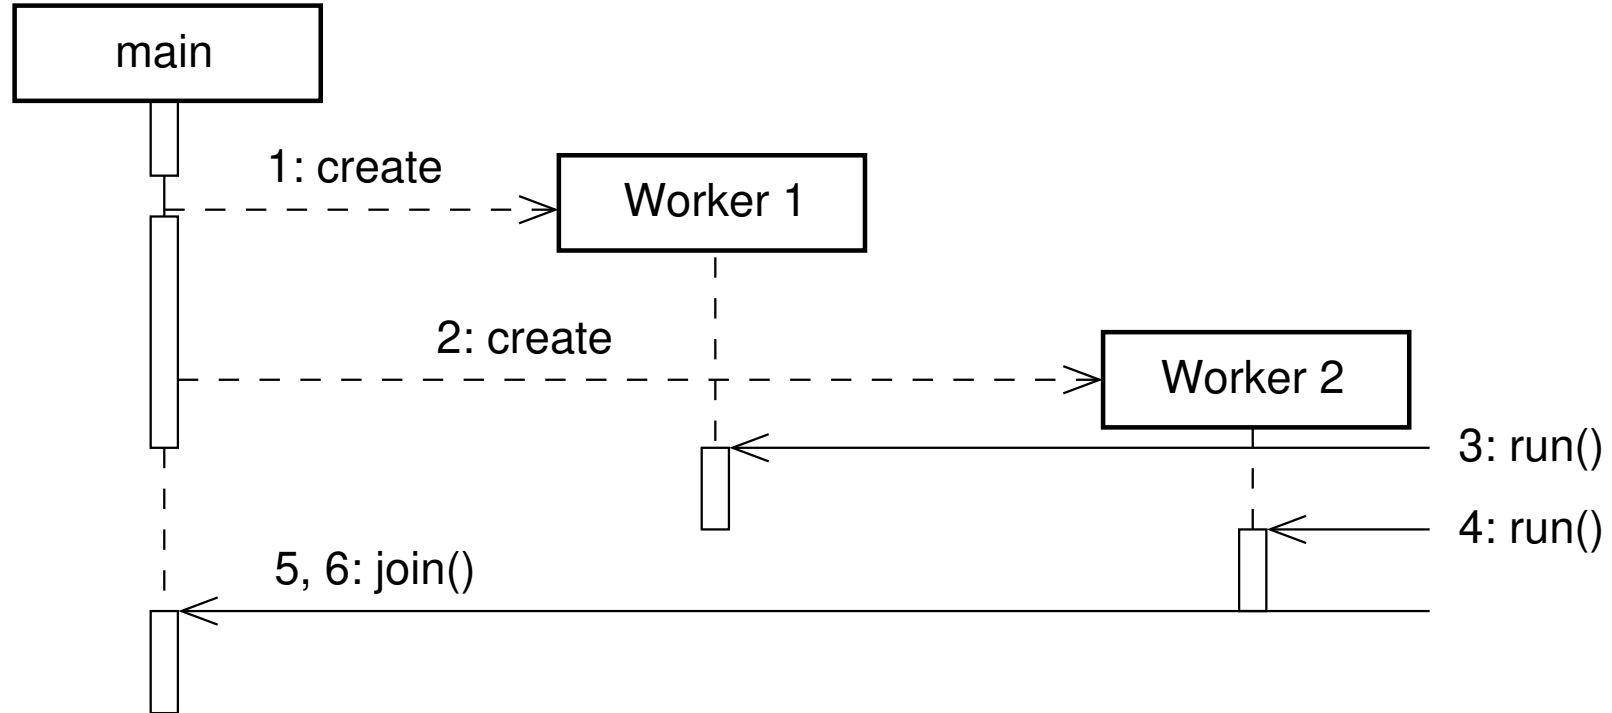
Concurrency in Java



In Java, a child thread is started as follows:

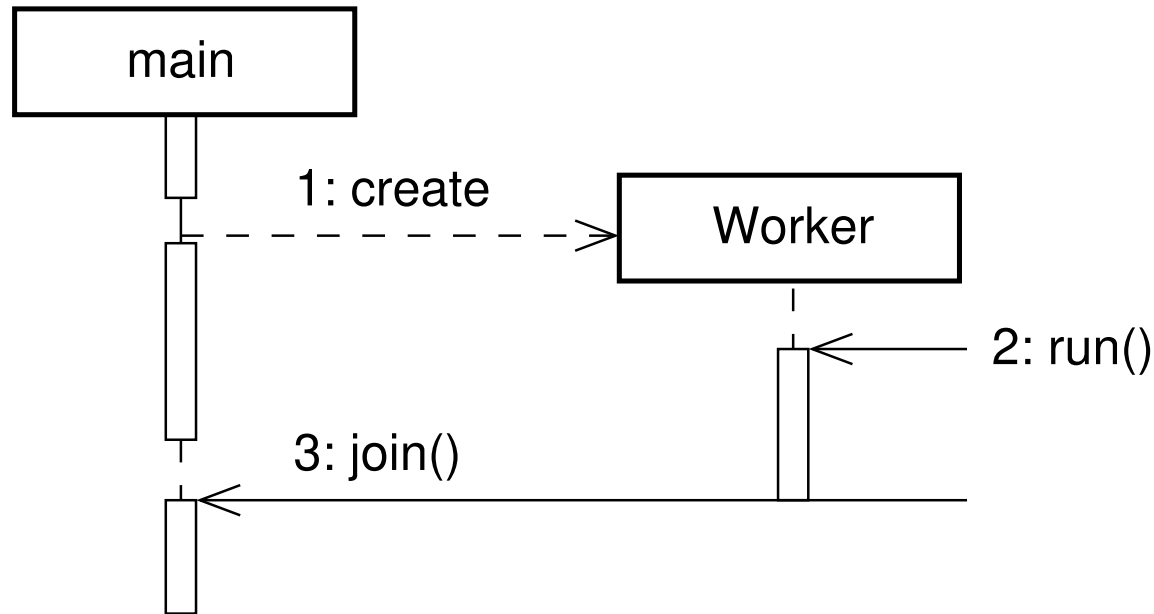
1. Create an instance of class **Thread**.
Alternative: Subclass of **Thread**, class implementing interface **Runnable**.
2. Call to **start** launches new thread as separate unit of execution.

Concurrency in Java (2)



- Constructor of **Thread** instance is initialized from main thread.
- „Payload” (function to be executed) of thread is in **run** method.
- Payload is executed from separate thread of execution („in parallel”).
- Execution of child thread may be immediate or delayed!

Fork/join parallelism



- **start** „spins off”/”forks” a new thread, executing its **run** method.
- Both main and child thread may run in parallel (given enough CPUs).
- **join** blocks current thread until target thread has terminated.

Fork/join by example

```
hw1.start(); // spin off first worker thread
hw2.start(); // spin off second worker thread
try {
    hw1.join(); // wait for first worker to finish
    hw2.join(); // wait for second worker
} catch (InterruptedException e) {
}
```

- Both child threads run independently.
- Parent waits for first child to terminate, then second one.
- **try/catch** block is just „boilerplate” here.
 - If current thread is interrupted while inside `join`, then exception is triggered.
 - This scenario cannot occur in this simple example.

The „payload” of HelloWorld

```
public class HelloWorld extends Thread {  
    StringBuffer buf;  
    String data;  
    HelloWorld(StringBuffer buf, String data) {  
        this.buf = buf;  
        this.data = data;  
    }  
    public void run() {  
        buf.append(data);  
    }  
}
```

- `run` appends a string to the (shared) buffer.
- In `main`:

```
hw1 = new HelloWorld(buf, "Hello, ");  
hw2 = new HelloWorld(buf, "World!");
```
- For expected result: `run` methods need to be executed in the right order.

Running HelloWorld in the normal JVM

```
% java HelloWorld  
Hello, World!
```

- Nothing special here.

Running HelloWorld in Java PathFinder

```
% jpf HelloWorld
JavaPathfinder v6.0 (rev 616) - (C) RIACS/NASA Ames Research Center
===== system under test
application: HelloWorld.java
===== search started: 11/7/11 1:42 PM
Hello, World!
Hello, World!
World!Hello,
Hello, World!
World!Hello,
===== results
no errors detected
===== statistics
elapsed time:          0:00:00
states:                new=68, visited=67, backtracked=134, end=5
search:                maxDepth=10, constraints=0
choice generators:     thread=68 (signal=0, lock=11, shared ref=16), data=0
heap:                  new=390, released=111, max live=352, gc-cycles=123
instructions:          4466
max memory:            81MB
loaded code:           classes=79, methods=1386
===== search finished: 11/7/11 1:42 PM
```

Explanation of the output

- Java PathFinder tries out all interleavings.
- Order in which both worker threads run is not restricted (enough)!
- **hw2** may run first, resulting in „World! Hello, ”.

This is interesting, but how did it happen?

- No properties specified → no errors found → no explanation.
- We need to add properties to our software!
- Simplest property: safety (assertion).

Automating the verification of the output

- Check value against expected value using assertion.
- Change line with print statement to
`assert (buf.toString().equals("Hello, World!"));`
- Now we compare the output to the expected value.
- Assertion failure results in program termination (error trace from JPF).

Error trace in JPF

```
~/jpf/jpf-core/bin/jpf HelloWorld.jpf > jpf.log
```

- We need the source path to see the source code statements in the trace.

```
===== error #1
gov.nasa.jpf.jvm.NoUncaughtExceptionsProperty
java.lang.AssertionError
  at HelloWorld.main(HelloWorld.java:25)
===== trace #1
----- transition #0 thread: 0
... // error trace shown here
===== results error #1:
gov.nasa.jpf.jvm.NoUncaughtExceptionsProperty
  "java.lang.AssertionError at HelloWorld.main(Hello..."
...
```

How to read the error trace

Thread number	Purpose	Thread name
0	main thread	main
1	add „Hello, ” to string buffer	Thread-0
2	add „World!” to string buffer	Thread-1

- Watch out for thread # changes in

```
----- transition #3 thread: 2
```

- Corresponding thread name of new thread is shown in

```
gov.nasa.jpfd.jvm.choice.ThreadChoiceFromSet... {Thread-1,>Thread-2}
```

- Number in thread name does not match thread number!
- Reading the trace, **write down summary of thread actions side by side.**

Fixing the example application

```
try {  
    hw1.join();  
} catch (InterruptedException e) {  
}  
hw2.start();  
try {  
    hw2.join();  
} catch (InterruptedException e) {  
}
```

- One possibility: create two **join** blocks.
- How to describe that change?

Highlighting the change: `diff`

- Automatically describes difference between text files.
- Applied to example source code:

```
% diff hw-assert/HelloWorld.java hw-fixed/HelloWorld.java
19d18
<  hw2.start();
21a21,24
>  } catch (InterruptedException e) {
>  }
>  hw2.start();
>  try {
```

- < refers to first file (old version), > to second file (new version).

Unified diff

```
% diff -u hw-assert/HelloWorld.java hw-fixed/HelloWorld.java
--- hw-assert/HelloWorld.java 2010-11-16 13:41:00.000000000 +0900
+++ hw-fixed/HelloWorld.java 2010-11-16 15:48:59.000000000 +0900
@@ -16,9 +16,12 @@
     HelloWorld hw1 = new HelloWorld(buf, "Hello, ");
     HelloWorld hw2 = new HelloWorld(buf, "World!");
     hw1.start();
-    hw2.start();
     try {
         hw1.join();
+    } catch (InterruptedException e) {
+    }
+    hw2.start();
+    try {
         hw2.join();
     } catch (InterruptedException e) {
     }
```

- – refers to old version, + to new version.
- Easier to read and more stable thanks to context (unchanged lines).

Mutual exclusion and locking

- Data race occurs because actions of thread are not atomic.
- Non-interference between actions can be ensured through „critical sections”.
 - Only one thread may enter a critical section at any time.
 - In Java, critical sections are implemented through **locking**.
 - Each critical section (on given data) has to use the same lock.
- **Locks:** „Gate keepers” to blocks of code (critical section).
 - Each lock may only be held by one thread at a given time.
 - If thread tries to obtain lock, then either...
 1. It atomically checks for its availability and obtains it, if possible, or
 2. It is suspended if the lock in question is not available.
 - Suspended thread may retry locking at any time.
- **Deadlock:**
 - Several threads are waiting for a resource that never becomes available.
 - „Deadly embrace”: Mutual dependency on resource owned by other threads.

Locking in Java: **synchronized**

- In Java, a thread may obtain a lock on any object instance.

```
Object lock = new Object(); // anything may be a lock
static class Lock {}
lock = new Lock();           // this is just „syntactic sugar“
```

- The keyword **synchronized** is used to obtain and release a lock:

```
synchronized (lock) {
    ... // critical section
}
```

- Lock is held inside the scope of **synchronized** block.
 - Lock is always obtained/release pairwise, within same method.
 - Lock can be obtain multiple times (*reentrancy*).
Reentrant operations only affect lock count.

More information about concurrency/locking in Java

Official tutorial:

<http://download.oracle.com/javase/tutorial/essential/concurrency/>

- Covers more detail than what is needed here.
- Recommended sections:
 - Defining and Starting a thread
 - Joins
 - Synchronization
 - Liveness
 - Guarded Blocks
- We will use some of the features covered there in future lectures.
Please study these sections before the first lab assignment.

Exercise 1, part 1: Dining Philosophers

```
public class DiningPhil {
    static class Fork {
    }
    static class Philosopher extends Thread {
        Fork left;
        Fork right;
        public Philosopher(Fork left, Fork right) {
            this.left = left;
            this.right = right;
            start();                // spin off new thread
        }
        public void run() {
            // think!
            synchronized (left) {
                synchronized (right) {
                    // eat!
                }
            }
        }
    }
}
```

Dining Philosophers: test harness

```
static final int N = 5;
public static void main(String[] args) {
    Fork[] forks = new Fork[N];
    for (int i = 0; i < N; i++) {
        forks[i] = new Fork();
    }
    for (int i = 0; i < N; i++) {
        new Philosopher(forks[i], forks[(i + 1) % N]);
    }
}
```

Dining Philosophers: Execution

- Normal JVM is unlikely to encounter deadlock:
`java DiningPhil`
should succeed, unlikely to encounter deadlock
- JPF finds possible deadlock:
`~/jpf/jpf-core/bin/jpf DiningPhil.jpf`

JPF output: part 1

```
JavaPathfinder v6.0 (rev 616) - (C) RIACS/NASA Ames Research Center
===== system under test
application: DiningPhil.java
===== error #1
gov.nasa.jpf.jvm.NotDeadlockedProperty
deadlock encountered:
  thread index=1,name=Thread-1,status=BLOCKED,
priority=5,lockCount=0,suspendCount=0
  thread index=2,name=Thread-2,status=BLOCKED,
priority=5,lockCount=0,suspendCount=0
  thread index=3,name=Thread-3,status=BLOCKED,
priority=5,lockCount=0,suspendCount=0
  thread index=4,name=Thread-4,status=BLOCKED,
priority=5,lockCount=0,suspendCount=0
  thread index=5,name=Thread-5,status=BLOCKED,
priority=5,lockCount=0,suspendCount=0
```

- **lockCount/suspendCount** can be ignored (refers to reentrant lock usage).

JPF output: part 2

```
===== trace #1
----- transition #0 thread: 0
gov.nasa.jpjf.jvm.choice.ThreadChoiceFromSet {>main}
    [2895 insn w/o sources]
    DiningPhil.java:50      : Fork[] forks = new Fork[N];
    DiningPhil.java:51      : for (int i = 0; i < N; i++) {
    DiningPhil.java:52      : forks[i] = new Fork();
    DiningPhil.java:23      : static class Fork {
    ...
    DiningPhil.java:31      : public Philosopher(Fork left, Fork right) {
    [188 insn w/o sources]
    DiningPhil.java:32      : this.left = left;
    DiningPhil.java:33      : this.right = right;
    DiningPhil.java:34      : start();
    [1 insn w/o sources]
```

- Summary of this transition:
create fork/philosopher instances, start threads.

JPF output: part 3

```
===== snapshot #1
thread index=1,name=Thread-0,status=BLOCKED,
priority=5,lockCount=0,suspendCount=0
  owned locks:DiningPhil$Fork@144
  blocked on: DiningPhil$Fork@145
  call stack:
    at DiningPhil$Philosopher.run(DiningPhil.java:40)
thread index=2,name=Thread-1,status=BLOCKED,
priority=5,lockCount=0,suspendCount=0
  owned locks:DiningPhil$Fork@145
  blocked on: DiningPhil$Fork@146
  call stack:
    at DiningPhil$Philosopher.run(DiningPhil.java:40)
...
```

- After all 10 transitions, thread status of remaining 5 threads is shown.
- All threads are blocked (waiting on lock) in this case: **deadlock!**

JPF output: part 4

```
===== results
error #1: gov.nasa.jpf.jvm.NotDeadlockedProperty
  "deadlock encountered:  thread id=1,name=Thread-1,..."
===== statistics
elapsed time:          0:00:01
states:                new=101, visited=136, backtracked=224, end=11
search:                maxDepth=16, constraints=0
choice generators:     thread=100 (signal=0, lock=25, shared ref=0), data=0
heap:                  new=387, released=1065, max live=376, gc-cycles=237
instructions:          7652
max memory:            81MB
loaded code:           classes=83, methods=1322
===== search finished
```

Exercise: Dining Philosophers

1. Create summary information of all thread actions:

Thread name/ID Trans.	main 0	Thread-0 1	...	Thread-4 5
0	create fork/ phil. instances, launch threads			
1-2		obtain lock x try to obtain lock y		

Complete table (and make sure to set correct lock IDs for x , y).
Lock ID should uniquely identify each lock.

2. Create lock dependency graph (for remaining 5 threads).

Nodes: thread ID + locks held by thread.

Edge $a \xrightarrow{x} b$: Thread a wants to obtain lock x held by thread b .

3. Propose possible **fix** for **program**, create „patch” (diff output).

Revised version of program should pass through JPF with no errors.

Lab exercise

1. Summary of thread actions: plain text or PDF (print to PDF from spreadsheet), scan from paper.
2. Lock dependency graph: PDF or scan from paper.
3. Copy file before editing: `cp DiningPhil.java DiningPhil.java.orig`.
Submit diff output:
`diff -u DiningPhil.java.orig DiningPhil.java > DiningPhil.patch.`

Requirements for revised philosophers

- Each philosopher still uses two forks together.
- Lock ordering may be changed.
- Number of philosophers = number of forks = N .
- Additional (non-Fork) locks may be used (although not necessary).
- Number of philosophers/forks remains at $N = 5$
(JPF is able to handle resulting state space).
- No „magic numbers” (constants other than 0, 1, N).
 N is defined exactly once in the code (as in existing version).
Use $N - 1, \dots \% N$, etc.
Code should still work when N is changed.
- **Hint:** If you need more than 5–10 extra lines of code,
then your solution is probably more complicated than necessary.

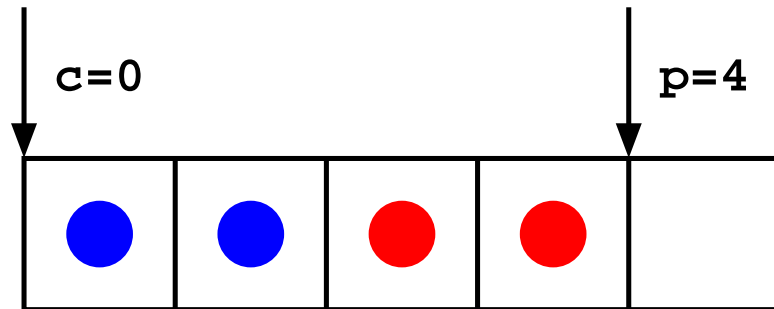
Shared conditions

- Locking can protect critical sections against concurrent access.
- Major building block for creating concurrent algorithms.
- Other operations required for more complex algorithms.

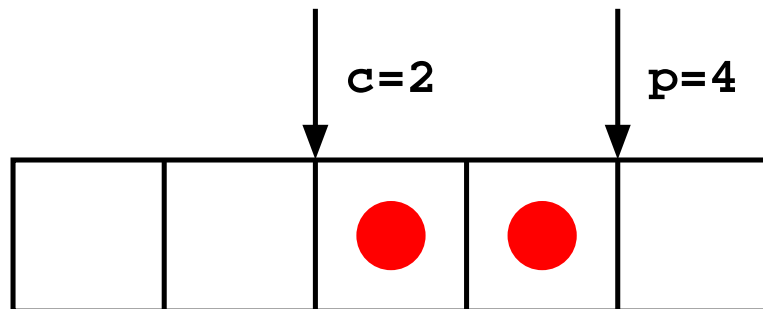
Producer/Consumer (Queue)

- Queue that allows concurrent data access to fixed-size buffer.
- Multi-element **put** blocks when queue is full (would overflow on put).
- Multi-element **remove** blocks when queue does not contain enough elements.
- Two indices:
 - Reader index **c** („consumer”).
 - Writer index **p** („producer”).
 - Both **c** and **p** may wrap around queue size (modulo).
 - Wrap-around: invariant $c + size \geq p$ (even after element insertion).


Queue implementation



After 2 put (2 elements each)



2 put, 1 remove



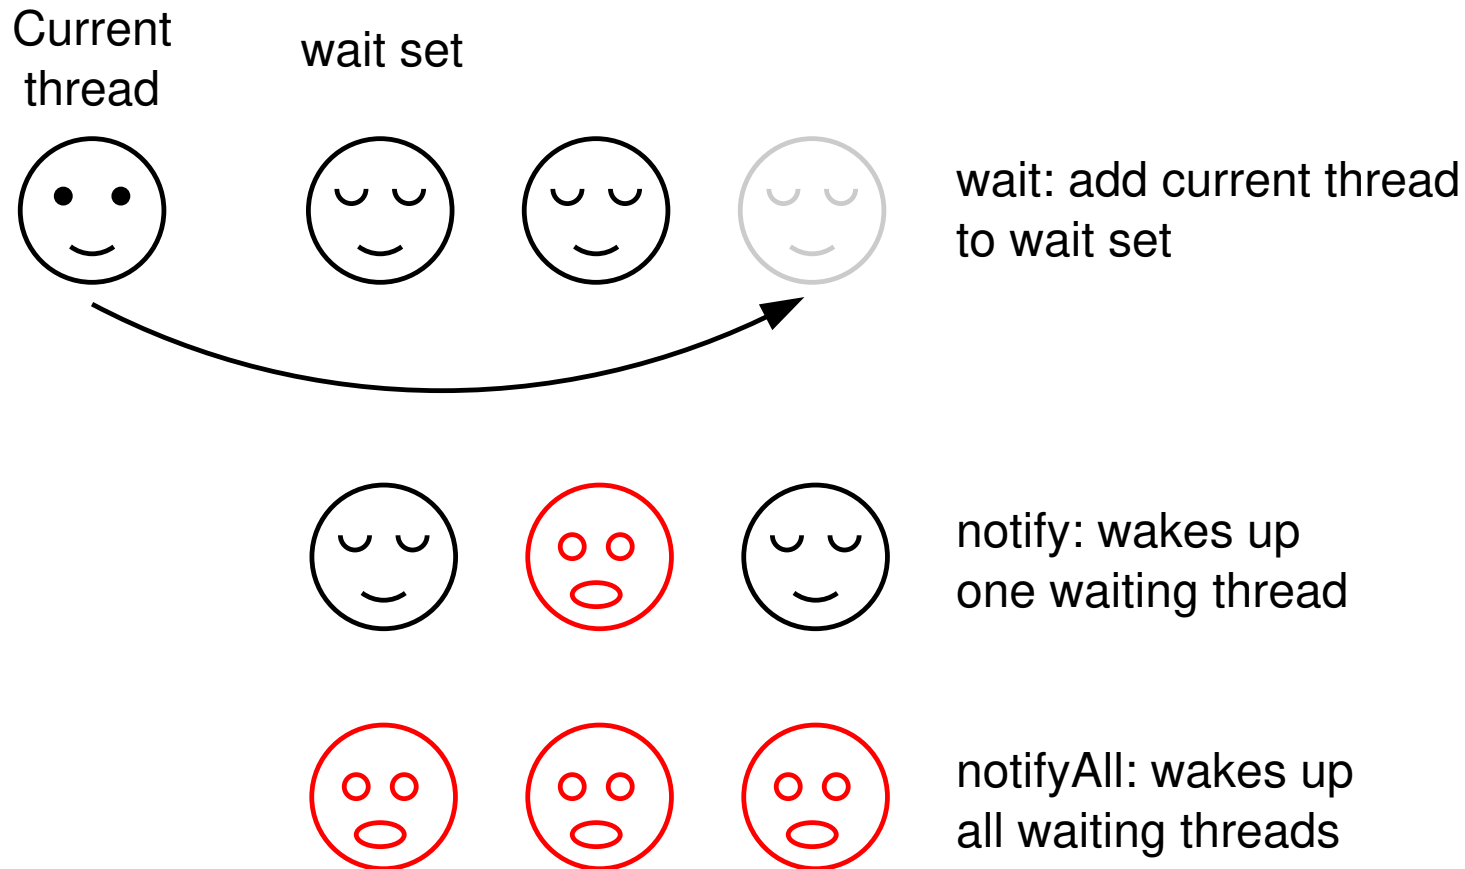
2 put, 1 remove, 1 put

Queue source code: remove

```
public synchronized void remove(byte[] storage) {
    int len = storage.length;
    waitForData(len);
    for (int i = 0; i < len; i++) {
        storage[i] = data[c++ % size];
    }
    /* index reduction to prevent index overflow */
    if (c >= size) {
        assert (p >= c);
        p -= size;
        c -= size;
    }
    notifyAll(); // wake up waiting producers
}
```

- **synchronized** ensures atomic behavior of operation.
- **waitForData** blocks until at least **len** elements are available.
- **notifyAll** signals any threads blocking on **put** (producers) that storage has increased after removal.

Shared conditions: `wait/notify`



- **`wait`** allows a thread to wait until another thread fulfills a (pre-)condition.
- **`notify/notifyAll`** signals waiting threads that condition is fulfilled.

wait/notify usage

- Lock on **x** must be held when **x.wait** or **x.notify** is called.
- **x.wait**:
 1. Releases lock on **x** (lock release not visible in source code).
 2. Puts current thread to sleep.
- **x.notify[All]**:
 1. Enables one or all threads to wake up.
 2. Lock on **x** is still held until **synchronized** block is exited!
 3. After release of lock **x**, newly awakened threads try to re-obtain lock **x** when scheduled.
- Even if all threads are woken up together, only one thread at a time, may take lock and continue!
- Condition waited on may no longer hold when awakened thread continues. This is a very frequent source of programming errors!

Waiting for something to happen

Assuming lock on `this` is held, *`await(condition)`* becomes six-liner in Java:

```
while (!condition) {  
    try {  
        wait();  
    } catch (InterruptedException e) {  
    }  
}
```

- Condition needs to be re-checked **before and after wait**.
- Monotonic conditions may be checked only before (using `if` instead of `while`).
- Inability of other threads to establish condition will result in deadlock (sometimes livelock).

Establishing a condition

Check condition

```
// wait for condition
synchronized (instance) {
    while (!condition) {
        try {
            instance.wait();
        } catch (InterruptedException e) {
        }
    }
    assert (condition);
    // holds unless data race present
}
```

Establish condition

```
// establish condition,
// notify waiting threads
synchronized (instance) {
    // always use locking when
    // modifying shared data
    condition = true;
    notifyAll();
}
```

- **condition** is usually not a boolean variable but an expression.
- Any data access that modifies part of **condition** (expression) needs synchronization!
- Synchronization may be done outside current method (by caller).

Queue source code: waitForData


```
private void waitForData(int len) {
    assert (len <= size);
    while (c + len > p) {
        try {
            wait();
        } catch (InterruptedException e) {
        }
    }
    assert (p >= c + len) :
        "p >= c + len failed: p = " + p + ", c = " + c + ", len = " + len;
    assert (p - size <= c) :
        "p - size <= c failed: p = " + p + ", c = " + c;
}
```

- **while** (!condition) \rightarrow condition holds at end of **while**/**wait** loop.
- $\neg(c + len > p) \rightarrow c + len \leq p \rightarrow p \geq c + len$
- Lock must be held outside this helper function;
then **waitForData** ensures availability of *len* elements.

Queue source code: put

```
/** Atomic put for multiple elements. Blocks until space available. */
public synchronized void put(byte[] items) {
    int len = items.length;
    assert (len <= size);
    // wait for space to become available
    while (p + len > c + size) {
        try {
            wait();
        } catch (InterruptedException e) {
        }
    }
    assert (p + len <= c + size) :
        "p + len <= c + size failed: p = " + p + ", c = " + c + ", len = " + len;
    for (int i = 0; i < len; i++) {
        data[p++ % size] = items[i];
    }
    assert (p <= c + size) :
        "p <= c + size failed: p = " + p + ", c = " + c;
    notifyAll(); // wake up waiting consumers
}
```

Producer/consumer dance



- Temporally, a producer must come first (as consumer would block).
- In program code, condition needs to be checked first.
- In Java, notifications *cannot* be tailored to producer or consumer threads (either all waiting threads, or one at random, no pre-determined subset).
- **notify** may wake up the „wrong” thread → use **notifyAll**, as we have threads waiting on *different* conditions.

Simple test bench: ProdCons

```
public class ProdCons {
    ... // declaration of constants
    static Queue q;
    static class Producer extends Thread {
        byte[] data;
        Producer(int i) {
            data = new byte[] { (byte)i, (byte)(i + DELTA) };
        }
        public void run() {
            q.put(data);
        }
    }
    ...
}
```

- Idea: test insertion of two elements at a time, $\langle i, i + \delta \rangle$.
- If retrieval of $\langle i, j \rangle$ is atomic, then $j = i + \delta$.

Test bench: Consumer and main

```
static class Consumer extends Thread {
    public void run() {
        byte[] result = new byte[2];
        q.remove(result);
        System.out.println(result[0] + ", " + result[1]);
        assert (result[1] - result[0] == DELTA);
    }
}

public static void main(String[] args) {
    q = new Queue(Q_SIZE);
    for (int i = 0; i < N; i++) {
        new Producer(i).start();
        new Consumer().start();
    }
}
```

- Result is printed on screen so JPF's state space search is visible.

Can we optimize the code?

- Obviously, at least one entry has to be produced before anything can be read.
- So any consumer thread has to wait until data is available.
- Should we always wait before checking the condition?
- JPF says no!

Thread number	Purpose	Thread name
0	main	main
1	Producer 1	Thread-1
2	Consumer 1	Thread-2
3	Producer 2	Thread-3
4	Consumer 2	Thread-4

Problem: Lost signal

- Data may already be available when producer thread comes first.
- No problem if more producers or consumers arrive and send redundant signal.
- Deadlock if
 1. No more producers (queue is full or no producers available).
 2. No more consumers (which would pass check and send signal).
- Signal from extra consumers would be redundant, mask possible deadlock.

Large-scale stress testing likely to produce redundant signals.

- Behavior may range from inefficiency to actual deadlock.
- Hard to diagnose with traditional testing.

Error trace for lost signal

Trans.	main 0	Thread-1 1	Thread-2 2	Thread-3 3	Thread-4 4
0–6	init. data, launch thr.				
7		add data notifyAll			
8			try to remove data wait		
9				try to add data wait	
10					try to remove data wait

- Signal is „lost” because consumer thread (#2) is too late to catch it.
- Checking the condition before calling `wait` resolves this problem.

Exercise1, part 2: wait/notify usage; test in ProdCons

- In **queue-notify**, `notifyAll` in `Queue.java` is replaced with `notify`.
- JPF now reports deadlock for `ProdCons` ($N = 2$, $Q_SIZE = 2$), 16 transitions.
jpf ProdCons.jpf

Hint:

- **Call to notify does not imply immediate signal delivery!**
- Signal delivery is shown by thread progressing past `wait` call:

```
----- transition #X thread: A
...
Queue.java:XX          : wait(); // DOES NOT GET SIGNAL
----- transition #Y thread: B
Queue.java:XX          : notify();
...                    // release lock, continue
----- transition #Z thread: C
Queue.java:XX          : wait();
Queue.java:XX          : }          // GETS SIGNAL
```


Tasks for lab exercise

1. **Analyze error trace, create summary information (table) for trace.**
Pay close attention to each `wait/notify` call, and received signals.
 - Which signal allows a thread to progress past the enclosing `while` loop?
 - Which signal is caught by a thread that *cannot* progress past the loop?
 - Would another thread be able to continue if it received that signal?
2. The error trace obviously shows that the program is incorrect.
Yet, another test program called `QueueTest` does not reveal any problem.
Study `QueueTest.java`; why can it not cause this particular failure?

Summary

- Simple testing may miss subtle concurrency defects.
- Java Pathfinder analyzes all possible thread schedules.
 - May uncover hard-to-find defects.
 - Requires properties to be verified (assertions).
- Error trace may be difficult to read.
 - Several threads involved, much detail.
 - Create summary by writing crucial thread actions side by side.